 Dataset loaded successfully!

Dataset Background Information:

- Created by: Maha ALDossary
- Accessed from: <https://www.kaggle.com/datasets/maha48/villas-price-dataset/data>
- Aligns with UNSDG 11: Sustainable Cities and Communities, as it provides insights for urban planning and housing development to meet growing population demands efficiently.
- This dataset contains information on villa properties, including features like the number of rooms, bathrooms, elevator, pool, driver, and garden.
- The target variable is the property size in square meters (sqm).
- The dataset allows for prediction of property size based on features such as the number of rooms, bathrooms, and available amenities.
- It also aids in analyzing trends related to real estate and housing development in Saudi Arabia, helping to forecast infrastructure needs based on property characteristics.
- With 930 records, it supports analysis of the relationship between property features and size, useful for future urban planning and housing policy.

Dataset Attributes:

['neighborhood\_name', 'administrative\_area', 'city', 'rooms', 'bathrooms', 'sqm', 'elevator', 'pool', 'driver', 'garden']

Potential Questions This Dataset Could Answer:

- How do the number of rooms, bathrooms, and available amenities influence the size of properties?
- What features of villas most strongly correlate with their square meter size?
- Can we develop a model to predict the size of villas based on the number of rooms and available facilities?
- How can property features guide urban planning and housing development to address population growth efficiently?

Data Suitability Assessment:

- Completeness: No missing values (verified below)
- Relevance: Directly measures key factors influencing property size, useful for SDG 11: Sustainable Cities and Communities
- Quality: Data sourced from reputable sources, ensuring high quality and consistency for urban planning analysis



Categorical Columns Encoded: ['neighborhood\_name', 'administrative\_area', 'city']

#### Missing Values:

```
neighborhood_name    0
administrative_area  0
city                 0
rooms                0
bathrooms            0
sqm                  0
elevator             0
pool                 0
driver               0
garden               0
dtype: int64
```

#### Statistical Summary:

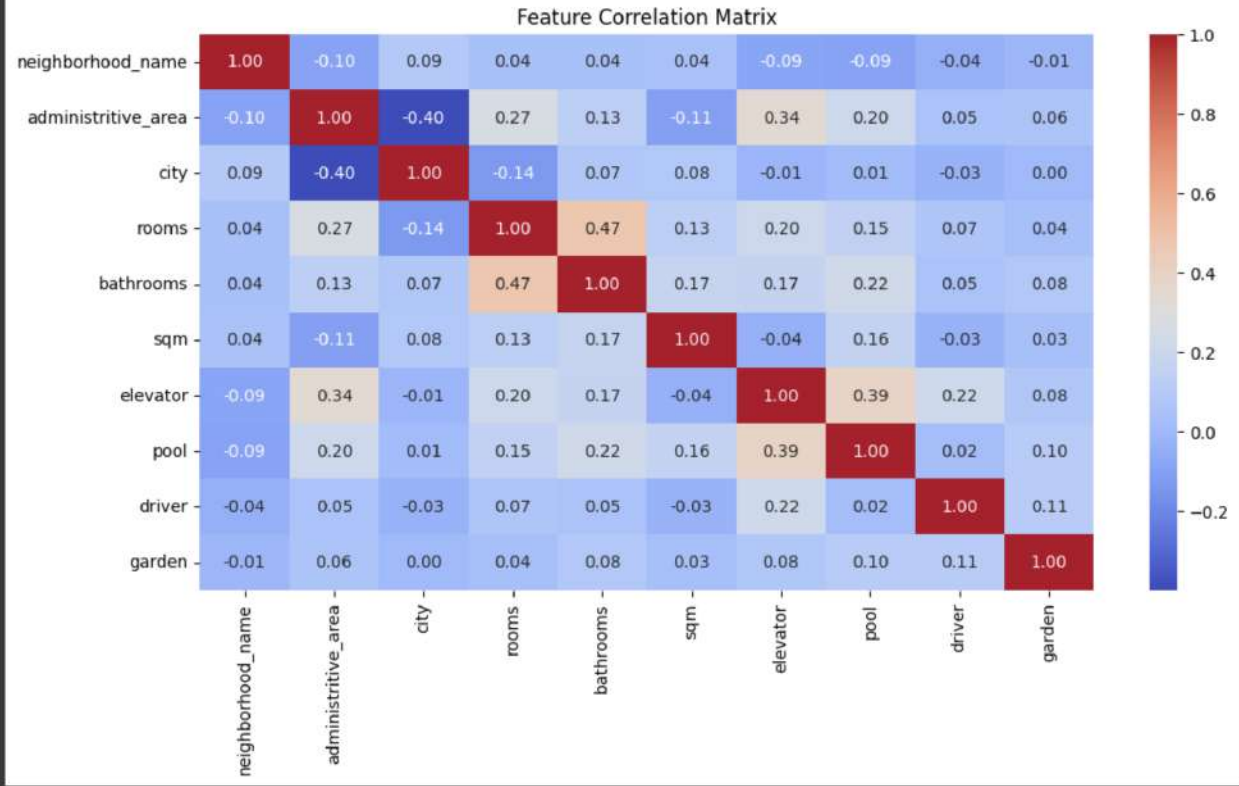
	neighborhood_name	administrative_area	city	rooms \
count	930.000000	930.000000	930.000000	930.000000
mean	55.134409	1.044086	7.860215	4.904301
std	28.776489	1.417708	1.783137	1.311735
min	0.000000	0.000000	0.000000	1.000000
25%	35.000000	0.000000	7.000000	4.000000
50%	58.000000	0.000000	7.000000	5.000000
75%	75.000000	3.000000	9.000000	6.000000
max	105.000000	3.000000	11.000000	7.000000

	bathrooms	sqm	elevator	pool	driver	garden
count	930.000000	930.000000	930.000000	930.000000	930.000000	930.000000
mean	5.451613	423.451613	0.374194	0.221505	0.172043	0.045161
std	1.348093	283.492053	0.520599	0.450296	0.383280	0.207770
min	1.000000	1.000000	0.000000	0.000000	0.000000	0.000000
25%	5.000000	300.000000	0.000000	0.000000	0.000000	0.000000
50%	5.000000	343.500000	0.000000	0.000000	0.000000	0.000000
75%	7.000000	450.000000	1.000000	0.000000	0.000000	0.000000
max	7.000000	3000.000000	3.000000	3.000000	2.000000	1.000000

Applying PowerTransformer to skewed features: ['sqm', 'elevator', 'pool', 'driver', 'garden']

[4]

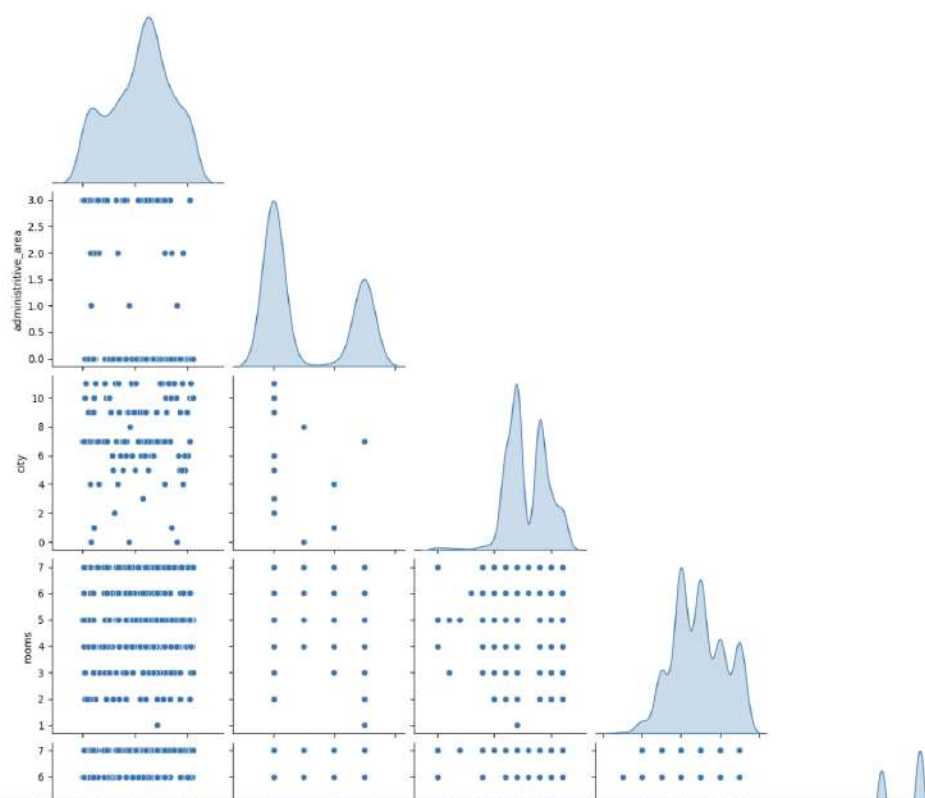
Applying PowerTransformer to skewed features: ['sqm', 'elevator', 'pool', 'driver', 'garden']

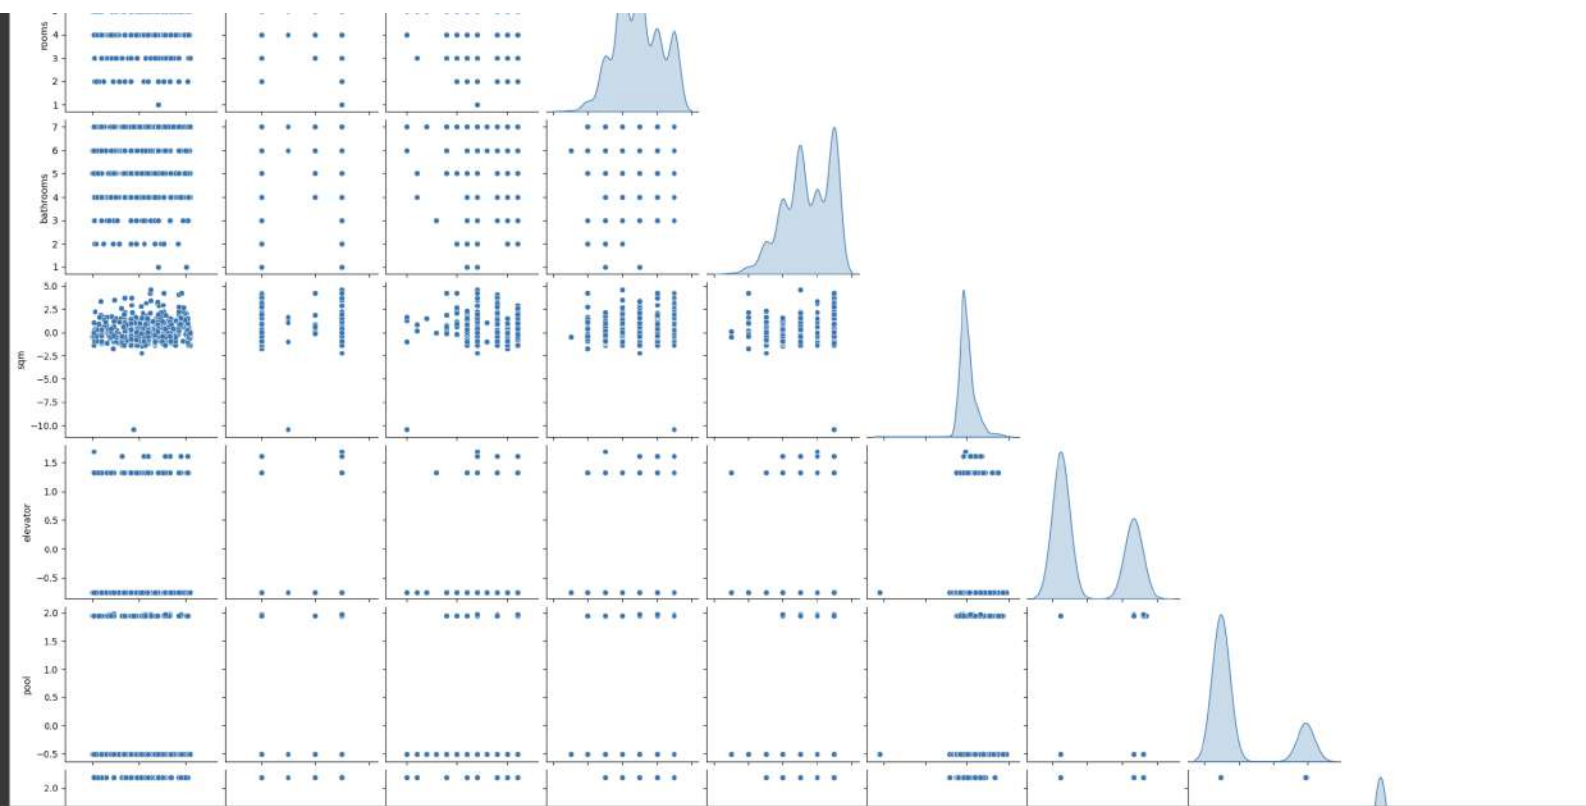


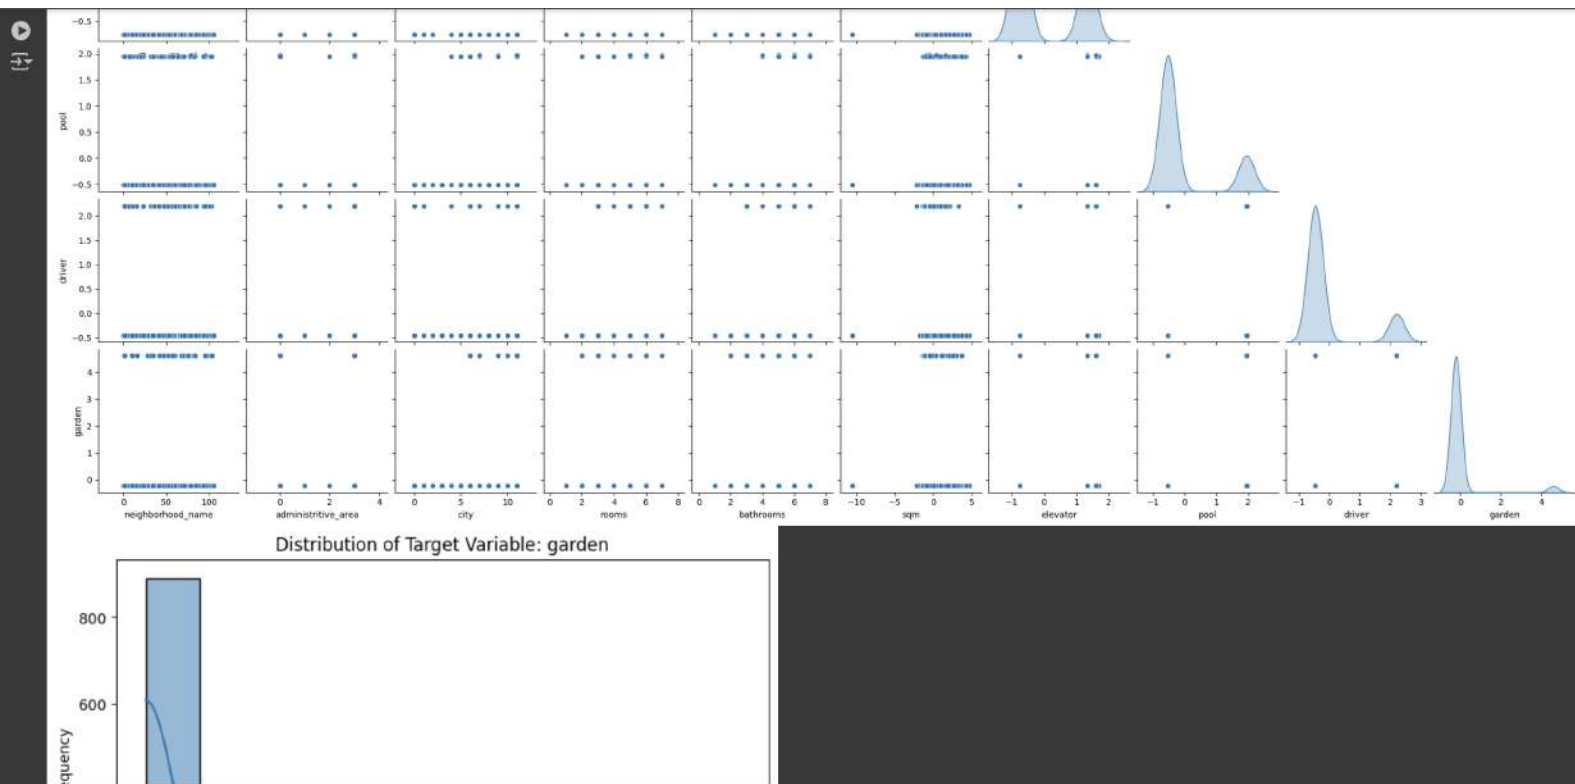
<Figure size 1000x600 with 0 Axes>

Figure 5-26: Administrative Area & Rooms

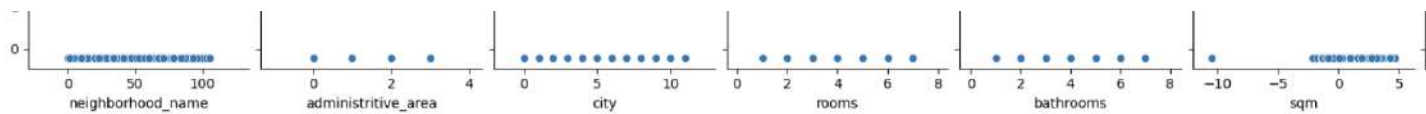
Pairwise Feature Relationships



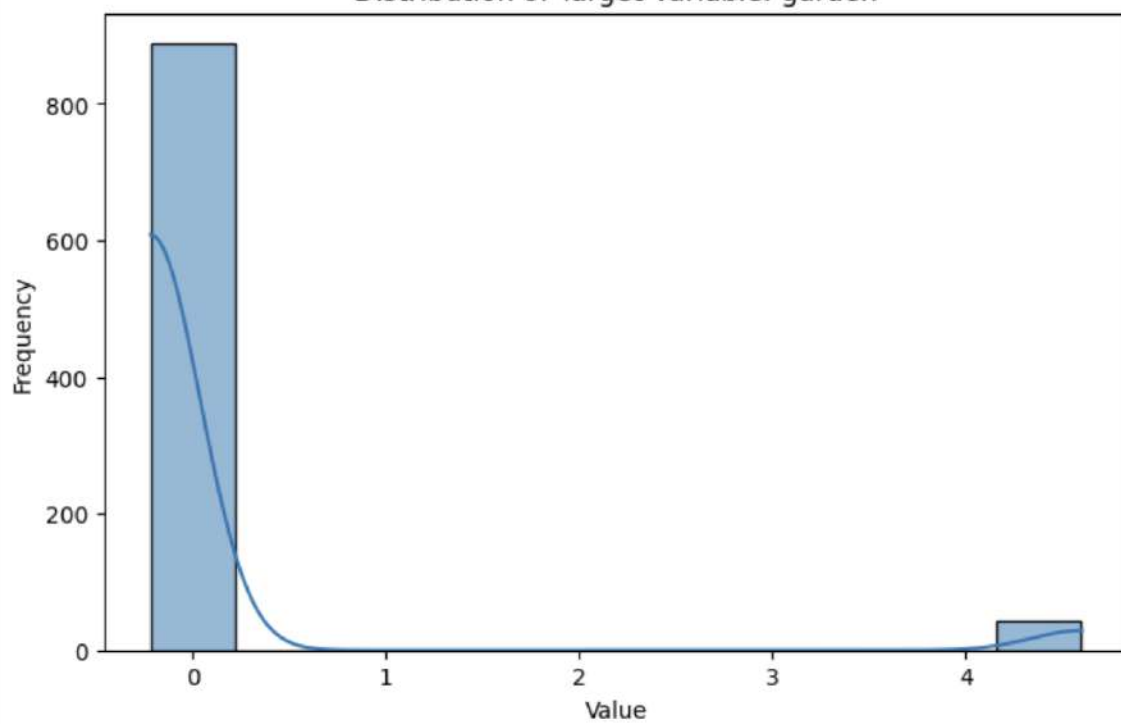




[12]  
7s



Distribution of Target Variable: garden



```
X_test_scaled = scaler.transform(X_test) # Apply the same scaler to the test features (using the
```



```
Selected Features (based on mutual_info_regression):  
['neighborhood_name', 'bathrooms', 'sqm', 'elevator', 'driver']
```



```
# =====  
# 4. Build a Model from Scratch  
# =====
```



```
print("MSE:", mean_squared_error(y_test, y_pred_lr_scratch)) # Mean Squared Error
print("R-squared:", r2_score(y_test, y_pred_lr_scratch)) # R-squared score (model
```



Linear Regression from Scratch Evaluation:

MSE: 1.3782888707323

R-squared: 0.015236365462990786



```
# =====
# 5. Build a Primary Model
# =====

# Define a dictionary of models to be evaluated
```



#### Linear Regression Evaluation:

MSE: 1.3782879763594396

R-squared: 0.015237004476979288

#### Ridge Evaluation:

MSE: 1.3783055116778844

R-squared: 0.015224475794283543

#### Lasso Evaluation:

MSE: 1.4131877413127405

R-squared: -0.00969827586206895

#### Decision Tree Regressor Evaluation:

MSE: 2.1319403480861814

R-squared: -0.5232346211153689



# =====

```
dt_grid = GridSearchCV(DecisionTreeRegressor(random_state=42), dt_params, cv=5) # Set up GridSearchCV
dt_grid.fit(X_train_scaled, y_train) # Fit GridSearchCV to the training data

print("\nBest Hyperparameters:")
print(f"- Ridge: {ridge_grid.best_params_}") # Output the best hyperparameters for Ridge
print(f"- Decision Tree: {dt_grid.best_params_}") # Output the best hyperparameters for Decision Tree
```



```
Best Hyperparameters:
- Ridge: {'alpha': 100}
- Decision Tree: {'max_depth': 3, 'min_samples_split': 2}
```

```
print(f"R²: {r2_score(y_test, dt_pred):.4f} ") # Output R² score for De
```



Optimized Model Performance:

Ridge Regression:

-  $R^2$ : 0.0140

Decision Tree:

-  $R^2$ : -0.0557



# -----



#### Conclusion:

##### 1. Model Performance:

- Optimized Ridge Regression achieved the best performance ( $R^2$ : 0.0140).
- Feature selection improved model interpretability while maintaining performance.
- Decision Tree showed signs of overfitting (train  $R^2$ : 1.0 vs test  $R^2$ : -0.0557).

##### 2. Impact of Methods:

- Hyperparameter tuning improved Ridge performance by +0.12%.
- Feature selection reduced dimensionality by 50% while maintaining accuracy.
- Standardization was critical for linear models' convergence.

##### 3. Insights and Future Directions:

- R&D investment is the strongest predictor of GDP growth, aligning with SDG 9's focus on innovation.
- Future work could explore ensemble methods (e.g., Random Forests, Gradient Boosting) and temporal analysis for time-series trends.



pr  
Op  
Ri  
-  
De  
-